

THE GEORGE WASHINGTON UNIVERSITY

WASHINGTON, DC

DATS 6313_10 Time Series Analysis & Modeling

Instructor: Dr. Reza Jafari

Time Series Final Project

Chiemeziem Oguayo

C.O May 2, 2022

Table of Contents

- **Abstract**
- **Introduction**
- **Descriptive Statistic/Dataset**
- **Stationary Check**
- **Time Series Decomposition**
- **Holt-Winter**
- **Feature Selection and Backwards Regression**
- **Multiple Linear Regression Model**
- **Base Modeling**
 - **Average Method**
 - **Naïve Method**
 - **Drift Method**
 - **SES Method**
- **ARMA Model**
- **LSTM Model**
- **Conclusion**
- **References**
- **Appendix**

Abstract

This project explores the use of time series analysis and modeling to identify a model that could predict an individual's household electric power consumption. First the time series data (electric power consumption), was interrogated to determine its various characteristics such, as trend, seasonal effects, stationarity, and descriptive statistics. Second, different time series modeling techniques were surveyed to model and forecast electric power consumption. These models were assessed using a battery of time series techniques. Ultimately, the best model found was to be "".

Introduction

The goal of the project was to apply the learning objectives taught in this course to a real dataset for modeling and prediction. The dataset that was obtained was from UCI machine learning repository. It is the measurement of electric power consumption of a home located in Sceaux, France. It is a multivariate time series data with 2075259 observations, that in turn could be used to model and forecast future household energy consumption. Here the target variable is the active energy, which is the energy consumed by the household.

Descriptive Statistics and Description of Dataset

As mentioned prior, the dataset is a multivariate time series data with 2075259 observations. It is a large dataset, with a huge computational load. So, the data was resampled over hours to reduce the load.

1. $(\text{global_active_power} * 1000 / 60 - \text{sub_metering_1} - \text{sub_metering_2} - \text{sub_metering_3})$ represents the active energy consumed every minute (in watt hour) in the household by electrical equipment not measured in sub-meterings 1, 2 and 3.
2. The dataset contains some missing values in the measurements (nearly 1,25% of the rows). All calendar timestamps are present in the dataset but for some timestamps, the measurement values are missing: a missing value is represented by the absence of value between two consecutive semi-colon attribute separators. For instance, the dataset shows missing values on April 28, 2007.

Attribute Information

- 1.date: Date in format dd/mm/yyyy
- 2.time: time in format hh:mm:ss
- 3.global_active_power: household global minute-averaged active power (in kilowatt)
- 4.global_reactive_power: household global minute-averaged reactive power (in kilowatt)
- 5.voltage: minute-averaged voltage (in volt)
- 6.global_intensity: household global minute-averaged current intensity (in ampere)
- 7.sub_metering_1: energy sub-metering No. 1 (in watt-hour of active energy). It corresponds to the kitchen, containing mainly a dishwasher, an oven and a microwave (hot plates are not

electric but gas powered).

8.sub_metering_2: energy sub-metering No. 2 (in watt-hour of active energy). It corresponds to the laundry room, containing a washing-machine, a tumble-drier, a refrigerator and a light.

9.sub_metering_3: energy sub-metering No. 3 (in watt-hour of active energy). It corresponds to an electric water-heater and an air-conditioner.




Figure 1. Plot of dependent variable vs time




Figure 2. ACF/PACF Plot of variable

From the ACF/PACF plot It seems that there is seasonality as there is a peak at every 12th lag, possibly indicating seasonality. Also, there is a cutoff at the PACF at lag 1, which may indicate a purely AR process.

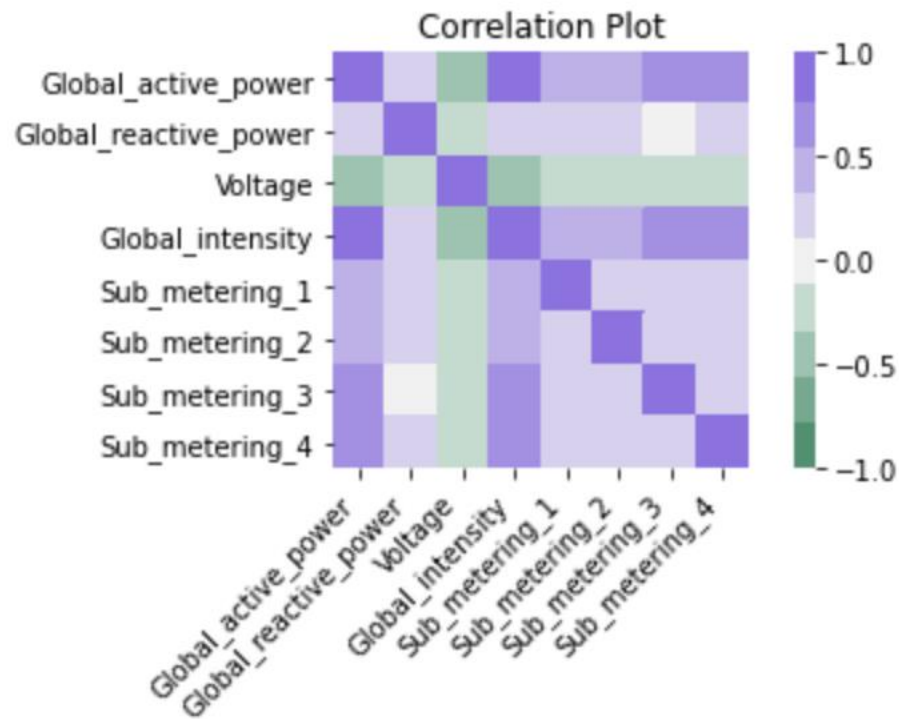


Figure 3. Correlation Plot of features

Global active power is strongly correlated with global intensity. And it seems that Global intensity shows similar correlations with the other variables as active power does with the other variables. In addition, all the variables are negatively correlated with voltage.

Stationarity

ADF Statistic: -14.263655

p-value: 0.000000

Critical Values:

1%: -3.431

5%: -2.862

10%: -2.567

p-value is less than 0.05, reject null hypothesis thus time series data is Stationary

Figure 4. ADF Test for stationarity

```
(0.6459031030714386,  
0.01,  
52,  
{'10%': 0.119, '5%': 0.146, '2.5%': 0.176, '1%': 0.216})
```

Figure 5. KPSS Test for stationarity

From the ADF and KPSS test we can determine that the time series is indeed stationary. As for the ADF test we rejected the null and for the KPSS test we failed to reject the null.

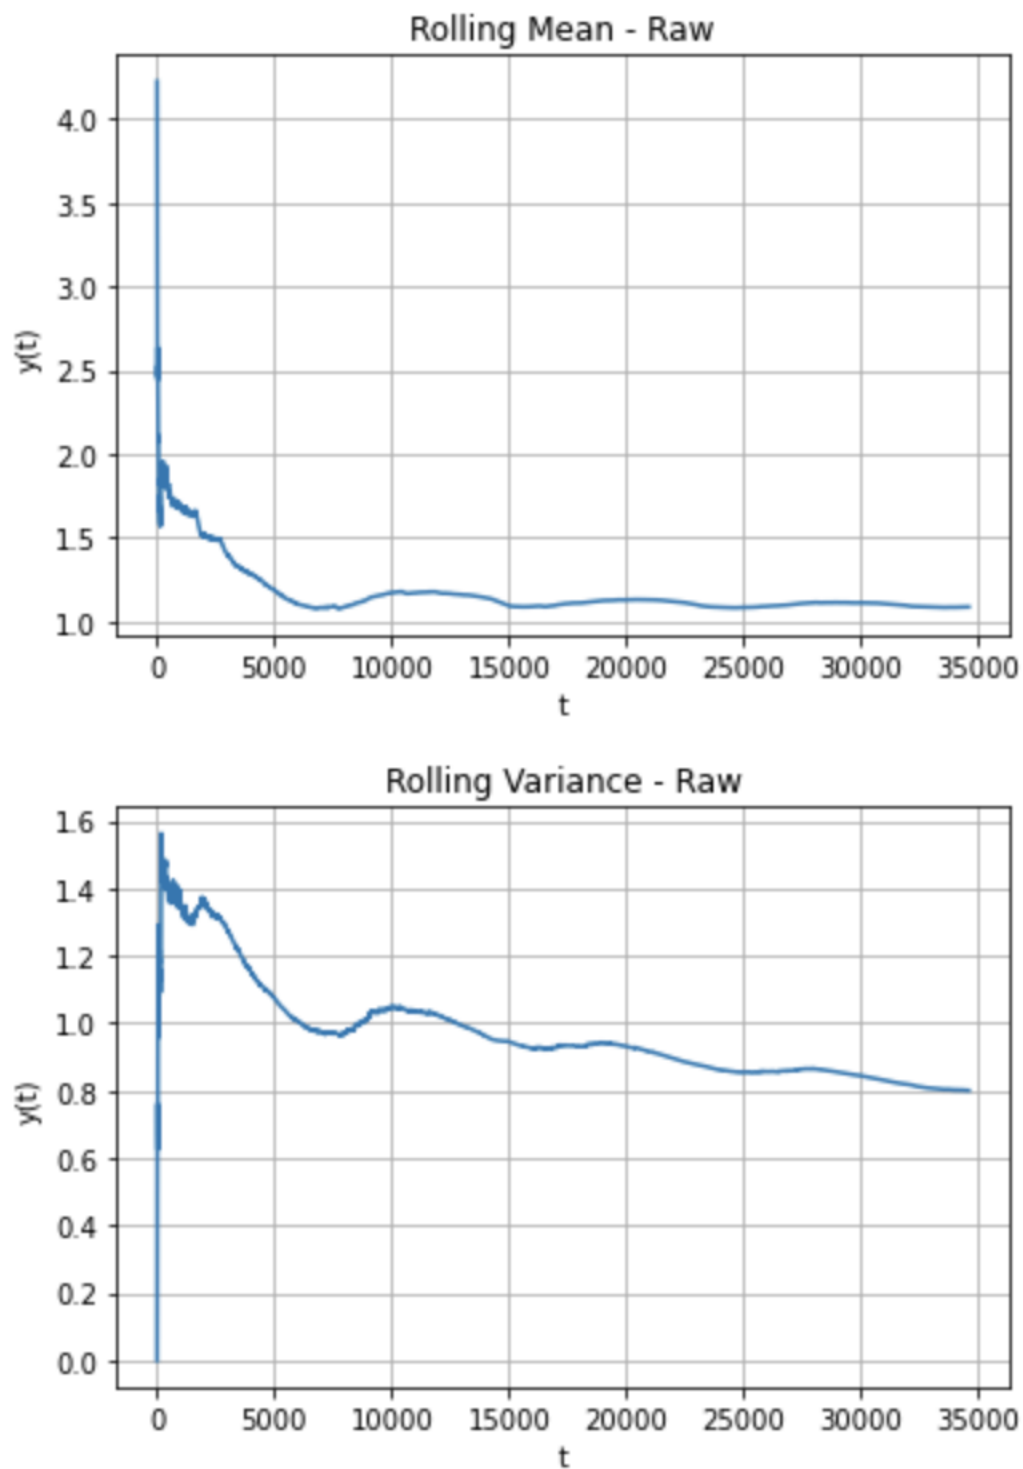


Figure 6. Plot of Rolling Mean and Variance

From the plot of the rolling mean and rolling variance we can see that both begin time stabilize as time progresses, further confirming that the time series is stationary.

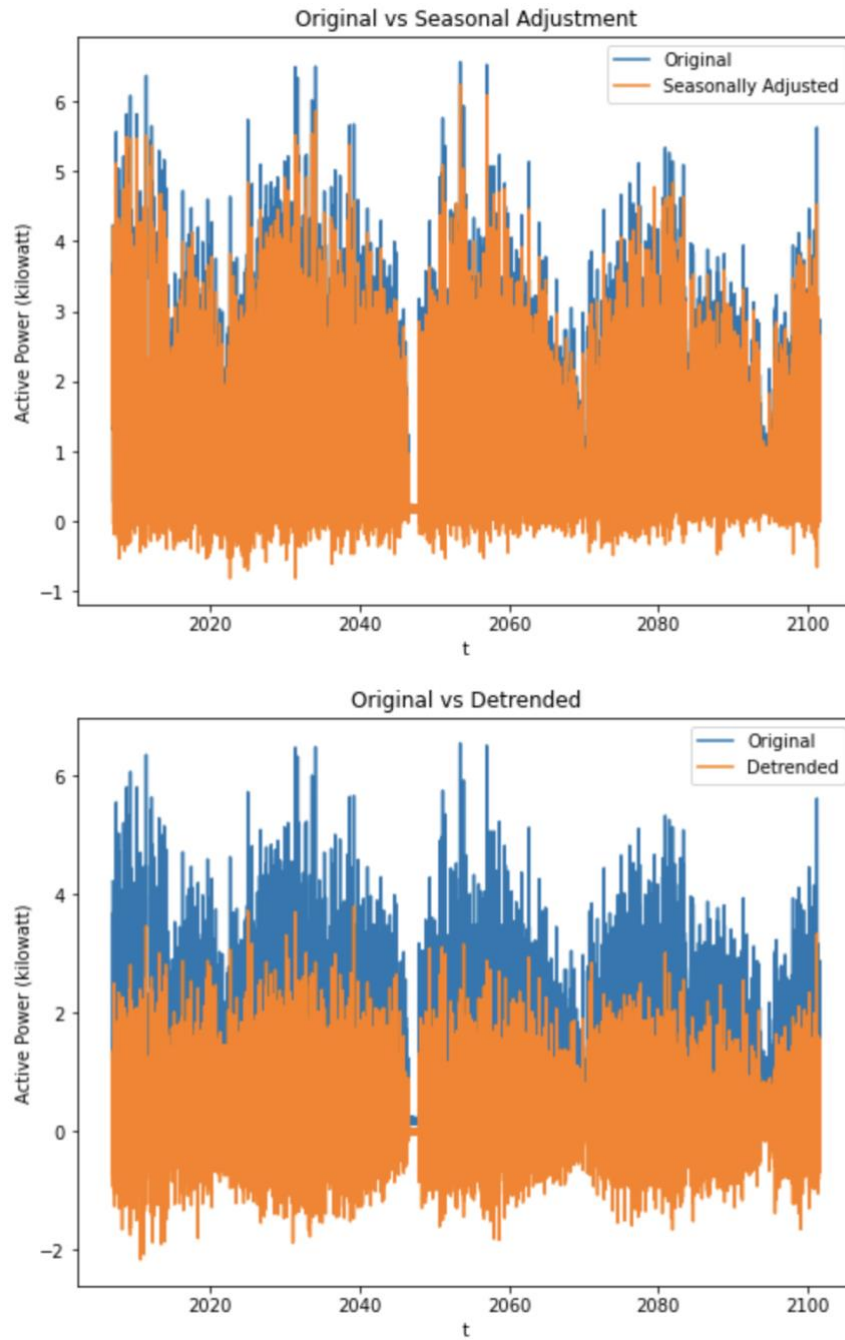


Figure 7. Plot of seasonally adjusted and detrended time series

The strength of trend for this dataset is 0.6678070388597215

The strength of seasonality for this dataset is 0.284315160227935

With the strength of the trend being 0.67 and strength of seasonality being 0.28, we can see that the trend in the data is detectable and it is more trended than it is seasonal.

Holt-Winters Forecast

The Holts-Winters forecasting method applies a triple exponential smoothing for trend, level, and seasonal components. Below are the results.




Figure 8. Holt-Winters Forecast

Model Evaluation

Plot of the time series data using Holt-Winters Forecast with seasonality of 12. Depicted is the train, test, and forecast of the test set.

The mean of the error of the HoltWinter Model is 0.2737996715713705

The variance of error of the HoltWinter Model is: 0.47283349535721647

The MSE of the HoltWinter Model is 0.5478

The RMSE of the HoltWinter Model is 0.7401351227985333




Figure 9. ACF of the residuals of the Holt-Winters Forecast

From the ACF plot we can see that the residuals are not white. Although, it decreases quickly, it seems that there are spikes at every 12th lag.

	lb_stat	lb_pvalue	bp_stat	bp_pvalue
20	5726.243605	0.0	5722.190184	0.0

From the Ljung box statistic (using 20 lags) we can reject the null and determine that model shows the lack of a good fit.

Feature Selection and Backwards Regression

Feature selection is the process of reducing the number of input variables when developing a predictive model. And backwards regression is starting with all the variables and at every step removing variables that do not add to the predictive power of the model.

OLS Regression Results						
Dep. Variable:	Global_active_power	R-squared:	1.000			
Model:	OLS	Adj. R-squared:	1.000			
Method:	Least Squares	F-statistic:	2.205e+30			
Date:	Mon, 02 May 2022	Prob (F-statistic):	0.00			
Time:	03:18:46	Log-Likelihood:	8.1482e+05			
No. Observations:	27671	AIC:	-1.630e+06			
Df Residuals:	27663	BIC:	-1.630e+06			
Df Model:	7					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
const	-2.753e-14	2.36e-14	-1.166	0.244	-7.38e-14	1.88e-14
Global_reactive_power	-7.55e-15	4.47e-15	-1.691	0.091	-1.63e-14	1.2e-15
Voltage	-3.469e-17	9.75e-17	-0.356	0.722	-2.26e-16	1.56e-16
Global_intensity	-6.661e-16	2.77e-15	-0.240	0.810	-6.1e-15	4.77e-15
Sub_metering_1	0.0600	7.09e-16	8.46e+13	0.000	0.060	0.060
Sub_metering_2	0.0600	7.08e-16	8.48e+13	0.000	0.060	0.060
Sub_metering_3	0.0600	6.75e-16	8.89e+13	0.000	0.060	0.060
Sub_metering_4	0.0600	6.93e-16	8.66e+13	0.000	0.060	0.060
Omnibus:	3045.788	Durbin-Watson:	0.002			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	4143.182			
Skew:	0.927	Prob(JB):	0.00			
Kurtosis:	3.397	Cond. No.	2.41e+04			

Figure 10. Regression results before selection

OLS Regression Results						
Dep. Variable:	Global_active_power	R-squared (uncentered):	1.000			
Model:	OLS	Adj. R-squared (uncentered):	1.000			
Method:	Least Squares	F-statistic:	3.271e+33			
Date:	Mon, 02 May 2022	Prob (F-statistic):	0.00			
Time:	03:18:52	Log-Likelihood:	9.0146e+05			
No. Observations:	27671	AIC:	-1.803e+06			
Df Residuals:	27665	BIC:	-1.803e+06			
Df Model:	6					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
Global_reactive_power	1.01e-14	1.7e-16	59.560	0.000	9.77e-15	1.04e-14
Voltage	-8.89e-18	1e-19	-88.899	0.000	-9.09e-18	-8.69e-18
Sub_metering_1	0.0600	3.15e-18	1.91e+16	0.000	0.060	0.060
Sub_metering_2	0.0600	2.49e-18	2.41e+16	0.000	0.060	0.060
Sub_metering_3	0.0600	1.47e-18	4.09e+16	0.000	0.060	0.060
Sub_metering_4	0.0600	1.26e-18	4.76e+16	0.000	0.060	0.060
Omnibus:	1248.722	Durbin-Watson:	0.321			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	4380.678			
Skew:	0.031	Prob(JB):	0.00			
Kurtosis:	4.948	Cond. No.	3.95e+03			

Figure 11. Regression Results after backwards selection

Backwards stepwise regression was the chosen method of reducing the number of features. At each step a variable was removed if the p-value was not significant or add to the model. In the end, the remaining variables were reactive power, voltage, Sub metering 1-4. All their p values were well under 0.05, so they added to the model. However, the R^2 did not differ between the full and reduced model. It remained at 1, however a model such as that may be meaningless.

The AIC and BIC were very low as well, suggesting this model is good a predicting energy consumption.

The condition number for X is: 3954.9106407340573

Singular Values are: [1.60768004e+09 2.37954293e+06 1.29161049e+06 5.05377498e+05 3.21781057e+05 1.02784178e+02]

The condition number 3954 is large, which may indicate that collinearity exist among the variables. No singular values were close to 0.

Multiple Linear Regression Model

The MLS was built from a continuation of the feature selection and backwards regression.




Figure 12. Multiple Linear Regression Model. The predictions/forecast overlaps very well with the train and test data

Model evaluation

The p-value of t-test is: Global_reactive_power 0.0

Voltage 0.0

Sub_metering_1 0.0

Sub_metering_2 0.0

Sub_metering_3 0.0

Sub_metering_4 0.0

The p-value of f-test is: 0.0

As this model uses the variables from the previous step, the variables in this model are all significant.

The AIC value of the model is: -1802906.2892009038

The BIC value of the model is: -1802856.9203596297

The R-squared value of the model is: 1.0

The Adjusted R-squared value of the model is: 1.0

The AIC and BIC are very small indicating that the model performs well at predicting energy Consumption. The R^2 and Adjusted R^2 are both 1 indicating a perfect fit, and has immense accuracy.

The variance of Prediction error is: 7.681051435276483e-31

The mean of Prediction error is: 1.479475695740853e-15

The variance of Forecast error is: 7.272602573587606e-31

The mean of Forecast error is: 1.3813634634340346e-15

The MSE of the residual is : 0.0

The Q value is: 4607.635496166898

We can see that the error for this is very small, with the MSE of this model being 0. This is a good model




Figure 13. ACF of MLR Forecast

We can say the error of the multiple linear regression model are white, as the ACF decays quickly.

Average Method




Figure 14. Average Method Forecast

Model evaluation




Figure 15. ACF of residuals of the residuals of Average Method Forecast

From the ACF we can see that the residuals are not white.

The MSE of the residual for the Average Method is: 0.55

The RMSE of the residual for the Average Method is: 0.7416198487095663

The variance of Forecast error is: 0.5340719263053189

The mean of Forecast error is: -0.12052942911981225

The Q value is: 5787.213651928438

This model is not good at capturing the variance in the data. Given the Q value the model also does not show a good fit.

Naïve Method




Figure 16. Naïve Method Forecast

Model evaluation




Figure 17. ACF of the residuals of Naïve Method Forecast

The MSE of the residual for the Average Method is: 0.93

The RMSE of the residual for the Average Method is: 0.9643650760992956
The variance of Forecast error is: 0.5340719263053189
The mean of Forecast error is: 0.6326776187722848
The Q value is: 5787.213651928438

We can see the error is higher in this model when compared to the previous models.

Drift Method




Figure 18. Drift Method Forecast




Figure 20. SES Method Forecast

Model Evaluation




Figure 21. ACF of the residuals of SES Forecast

The MSE of the residual for the Average Method is: 0.73

The RMSE of the residual for the Average Method is: 0.8544003745317531

The variance of Forecast error is: 0.5340719263053189

The mean of Forecast error is: 0.4377268455074007

The Q value is: 5787.213651928447

This model is not good at capturing the variance in the data. Given the Q value the model also does not show a good fit. As with some of the previous models, there still appears to be a seasonal trend (except MLR).

Arima Model (1,0)

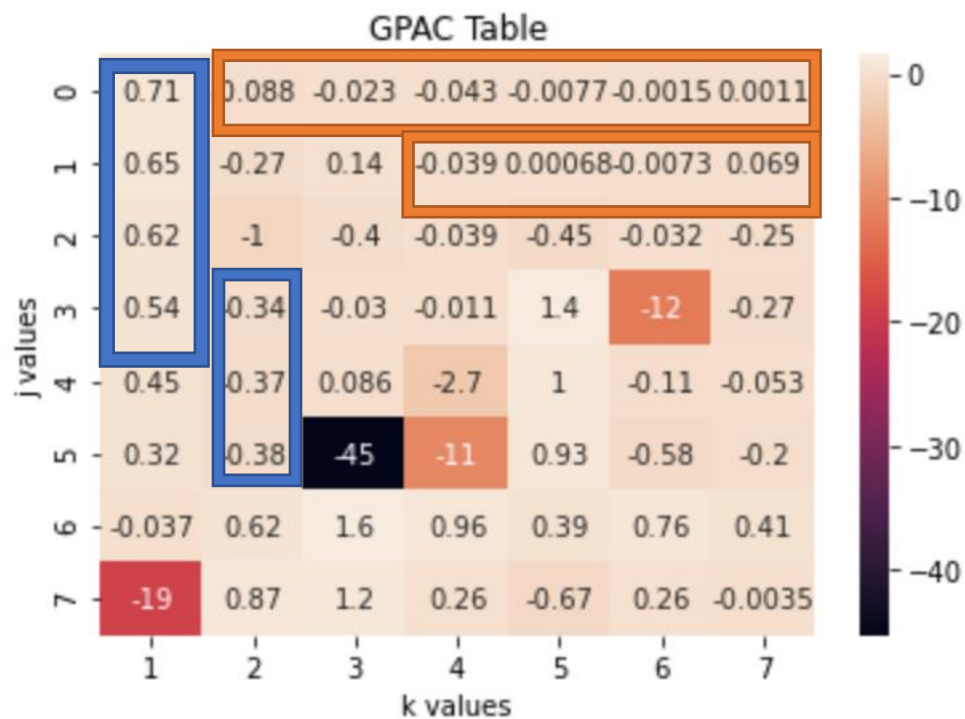


Figure 22. GPAC used to determine order of ARMA

Potential orders of (1,0), (2,1), (2,0)

ARMA Model Results						
Dep. Variable:	Global_active_power	No. Observations:	34589			
Model:	ARMA(1, 0)	Log Likelihood	-34488.616			
Method:	css-mle	S.D. of innovations	0.656			
Date:	Mon, 02 May 2022	AIC	68981.233			
Time:	03:21:13	BIC	68998.135			
Sample:	12-16-2006	HQIC	68986.620			
	- 11-26-2010					
	coef	std err	z	P> z	[0.025	0.975]
ar.L1.Global_active_power	0.8849	0.003	353.162	0.000	0.880	0.890
ROOTS						
	Real	Imaginary	Modulus	Frequency		
AR.1	1.1301	+0.0000j	1.1301	0.0000		

Figure 23. Arima Model (1,0) Summary

The parameters(coef) and confidence interval is highlighted in the box above.

The AR estimated coefficient a_0 is: 0.8848861970081253

The confidence interval for estimated coefficient a_0 is:

ar.L1.Global_active_power 0.879975 0.88979

The confidence interval for the estimated parameter does not include 0 so they are significant.

No zero/pole cancellations




Figure 23. ARMA Model (1,0)

It seems that a good amount of the test data is captured by the model




Figure 24. ACF of residuals of ARMA (1,0)

Model evaluation

The standard deviation of the parameter estimates is: 0.1849904247716646

The MSE of the residual for the Average Method is: 0.02

The RMSE of the residual for the Average Method is: 0.1414213562373095

The variance of Prediction error is: 0.4397077869371758

The mean of Prediction error is: 0.12782308641163162

The variance of Forecast error is: 0.007077087308328619

The mean of Forecast error is: 0.11393706581311966

The variance of prediction vs forecast error is: 62.1311802130389

The Q value is: 5787.213651928447

The covariance of estimated parameters is:

ar.L1.Global_active_power 0.000006

This model has a low MSE compared to other models and seems to perform well. However, the ACF of the residuals are not white. The seasonal trends seem to be apparent in the lags. This is not the best model.

ARMA (2,1)

ARMA Model Results

Dep. Variable:	Global_active_power	No. Observations:	34589
Model:	ARMA(2, 1)	Log Likelihood	-32586.234
Method:	css-mle	S.D. of innovations	0.621
Date:	Tue, 03 May 2022	AIC	65180.468
Time:	03:48:47	BIC	65214.273
Sample:	12-16-2006	HQIC	65191.242
	- 11-26-2010		

	coef	std err	z	P> z	[0.025	0.975]
ar.L1.Global_active_power	1.6807	1.77e-05	9.51e+04	0.000	1.681	1.681
ar.L2.Global_active_power	-0.6807	1.34e-05	-5.08e+04	0.000	-0.681	-0.681
ma.L1.Global_active_power	-0.9954	0.001	-1902.583	0.000	-0.996	-0.994

	Real	Imaginary	Modulus	Frequency
AR.1	1.0000	+0.0000j	1.0000	0.0000
AR.2	1.4691	+0.0000j	1.4691	0.0000
MA.1	1.0047	+0.0000j	1.0047	0.0000

Figure 25. ARMA (2,1) Model Results




Figure 26. ARMA (2,1) Forecast

The AR coefficient a_0 is: 1.680670545689072

The AR coefficient a_1 is: -0.680677527114449

The MA coefficient a_0 is: -0.9953711208984621

The confidence interval for estimated coefficient is:

```
ar.L1.Global_active_power 1.680636 1.680705
ar.L2.Global_active_power -0.680704 -0.680651
ma.L1.Global_active_power -0.996397 -0.994346
```

the covariance Matrix for the data is

```
ar.L1.Global_active_power \
ar.L1.Global_active_power      3.124721e-10
ar.L2.Global_active_power      -1.810637e-10
ma.L1.Global_active_power      -1.752500e-09
```

```
ar.L2.Global_active_power \
ar.L1.Global_active_power      -1.810637e-10
ar.L2.Global_active_power      1.793209e-10
ma.L1.Global_active_power      2.802235e-11
```

```
ma.L1.Global_active_power
ar.L1.Global_active_power      -1.752500e-09
ar.L2.Global_active_power      2.802235e-11
ma.L1.Global_active_power      2.737050e-07
```

Model Evaluation




Figure 27. ACF of residuals of ARMA(2,1)

The MSE of the residual for the Average Method is: 0.05

The RMSE of the residual for the Average Method is: 0.22360679774997896

The variance of Prediction error is: 0.4099483755224675

The mean of Prediction error is: 0.0006077393701143024

The variance of Forecast error is: 0.04926784273669454

The mean of Forecast error is: -0.0018049658070193994

The variance of prediction vs forecast error is: 8.320810345063865

The Q value is: 5032.4482524419

The standard deviation of the parameter estimates is: 0.14843194803134716

This model has a low MSE compared to other models and seems to perform well. However, the ACF of the residuals are not white. The seasonal trends seem to be apparent in the lags. This is not the best model. And the ARMA (1,0) performs better than with order (2,1)

LSTM

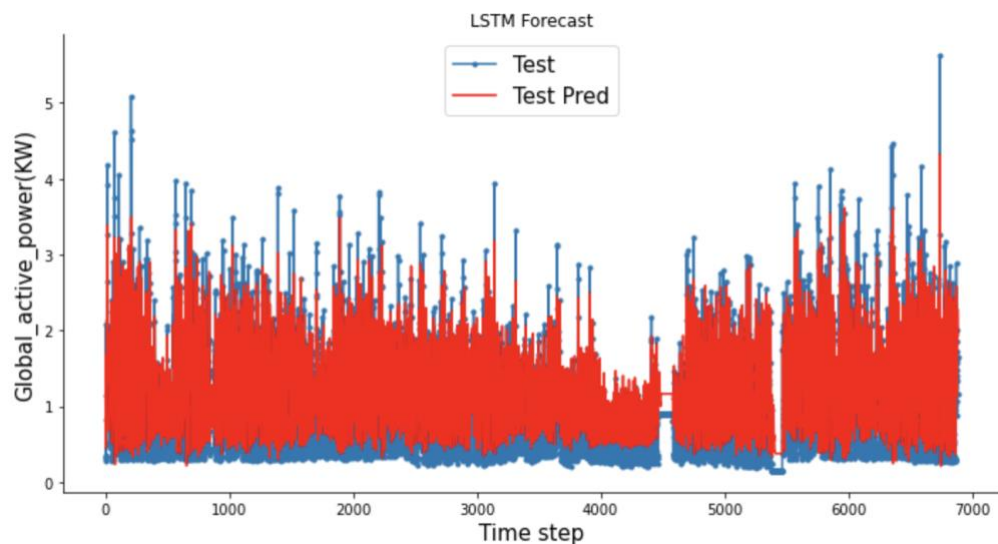


Figure 28. LSTM Forecast




Figure 29. Model loss. The loss of the model is very low for the test set

Train Mean Absolute Error: 0.4644396976856607

Train Root Mean Squared Error: 0.6176469288002417

Test Mean Absolute Error: 0.4240039845240055

Test Root Mean Squared Error: 0.5469491639580324

The Q value is: 79.6109408324576




Figure 30. ACF of residuals of LSTM forecast

The ACF exhibits a spike at 0, with most of the other spikes in the insignificant region.

```
lb_stat lb_pvalue bp_stat bp_pvalue
20 79.73596 4.352548e-09 79.610941 4.570273e-09
```

When interpreting the Q value and the results of the Ljung box test, we can see that the model shows a good fit.

Conclusion

Model	MSE
Holt-Winter's Model	0.74
Multiple Linear Regression Model	0.0
Average	0.55
Naïve	0.93
Drift	1.84
SES	0.73
ARMA (0,1)	0.02
ARMA (2,1)	0.05
LSTM	0.55

The final model that will be chosen is the Multiple Linear Regression Model, as it had the lowest MSE and was the only model with a perfect fit. This model was better able to capture the data. Although it should be noted that the ARMA models performed just as well, with MSE's almost comparable to that of the Multiple Linear Regression model. The difference in MSE's could be said to be negligible. Also, interesting to note, is that the LSTM model, was the only

model whose Q value was significant. The final forecast function can be written using the weights present in Multiple Linear Regression summary.

$$Y = (1.01e-14) * (\text{global_reactive_power}) - (8.98e-18) * (\text{Voltage}) + (0.06) * (\text{Sum (Sub meterings 1-4)})$$




Figure 31. H step prediction of best model – Multiple Linear Regression Model

Reference:

Dataset source:

<https://archive.ics.uci.edu/ml/datasets/Individual+household+electric+power+consumption#>

Li, Susan. "Time Series Analysis, Visualization & Forecasting with LSTM." *Medium*, Towards Data Science, 17 May 2019, <https://towardsdatascience.com/time-series-analysis-visualization-forecasting-with-lstm-77a905180eba>.

Appendix

***See pynb file attached if you want to run the code.**

```
#!/usr/bin/env python
```

```
# coding: utf-8
```

```
# In[164]:
```

```
import matplotlib.pyplot as plt
```

```
import numpy as np
```

```
import pandas as pd
```

```
import seaborn as sns
```

```
import statsmodels.api as sm
```

```
import statsmodels.tsa.holtwinters as ets
```

```
from statsmodels.tsa.api import SimpleExpSmoothing
```

```
from scipy import signal
```

```
import math
```

```
import seaborn as sns
```

```
from statsmodels.graphics.tsaplots import plot_acf , plot_pacf
```

```
from statsmodels.tsa.stattools import adfuller
```

```

from statsmodels.tsa.stattools import acf

from sklearn.model_selection import train_test_split

from numpy import linalg as LA

from datetime import datetime

```

```

# # READING DATA AND CREATING SEVERAL TIME SERIES FUNCTIONS

```

```

# In[165]:

```

```

df = pd.read_csv('household_power_consumption.txt', sep=';',
                 parse_dates={'dt': ['Date', 'Time']}, infer_datetime_format=True,
                 low_memory=False, na_values=['nan', '?'], index_col='dt')

```

```

# In[166]:

```

```

# Creating a new columns for all the engery in the rest of the house that doesn't include
submeterings 1,2,&3.

```

```

# We are only given sub_metering1,sub_metering2,sub_metering3.

```



```
# (global_active_power*1000/60 - sub_metering_1 - sub_metering_2 - sub_metering_3)
```

represents the active energy consumed

```
# every minute (in watt hour) in the household by electrical equipment not measured in sub-meterings 1, 2 and 3.
```

```
df['Sub_metering_4'] = (
    df['Global_active_power'] * 1000 / 60 - df['Sub_metering_1'] - df['Sub_metering_2'] -
    df['Sub_metering_3'])
```

```
# In[167]:
```

```
def difference(dataset, interval):
    diff = []
    for i in range(interval, len(dataset)):
        value = dataset[i] - dataset[i - interval]
        diff.append(value)
    return diff

def ADF_Cal(x):
    result = adfuller(x)
    print("ADF Statistic: %f" % result[0])
    print('p-value: %f' % result[1])
```

```

print('Critical Values:')

for key, value in result[4].items():

    print('\t%s: %.3f' % (key, value))

if result[1] < 0.05:

    print("p-value is less than 0.05, reject null hypothesis thus time series data is Stationary")

else:

    print("p-value is greater than 0.05, we failed to reject null hypothesis thus time series data
is "

        "Non-Stationary")

```

```

from statsmodels.tsa.stattools import kpss

```

```

def kpss_test(timeseries):

    print('Results of KPSS Test:')

    kpsstest = kpss(timeseries, regression='c', nlags="auto")

    kpss_output = pd.Series(kpsstest[0:3], index=['Test Statistic', 'p-value', 'Lags Used'])

    for key, value in kpsstest[3].items():

        kpss_output['Critical Value (%s)' % key] = value

    print(kpss_output)

def rolling_mean(x):

```

```

roll = []

for i in range(1, len(x) + 1):

    rolling = np.mean(x[:i])

    roll.append(rolling)

return roll


def rolling_var(x):

    roll = []

    for i in range(1, len(x) + 1):

        rolling = np.var(x[:i])

        roll.append(rolling)

    return roll


def acf_cal(x, k):

    acf_values = []

    mn = np.mean(x)

    for k in range(0, k + 1):

        acf_values.append(sum((x - mn).iloc[k:] * (x.shift(k) - mn).iloc[k:]) / sum((x - mn) ** 2))

    return acf_values


def GPAC(ry, j0, k0):

    # get phi

    def phi(ry, j, k):

```

```

# FIRST STEP: GET phi

# creating 0 for placeholders for denominator
denominator = np.zeros(shape=(k, k))

# replacing denom matrix with ry(j) values
for a in range(k):
    for b in range(k):
        denominator[a][b] = ry[abs(j + a - b)]

# making a copy of denom for numerator
numerator = denominator.copy()

# creating last column for numerator
numL = np.array(ry[j + 1:j + k + 1])
numerator[:, -1] = numL

phi = np.linalg.det(numerator) / np.linalg.det(denominator)

return phi

table0 = [[0 for i in range(1, k0)] for i in range(j0)]

for c in range(j0):
    for d in range(1, k0):
        table0[c][d - 1] = phi(ry, c, d)

pac = pd.DataFrame(np.array(table0), index=np.arange(j0), columns=np.arange(1, k0))

```

```
return pac
```

```
def GPAC_plot(acf, j, k):  
  
    gpac = GPAC(acf, j, k)  
  
    plt.figure()  
  
    sns.heatmap(gpac, annot=True)  
  
    plt.title('GPAC Table')  
  
    plt.xlabel('k values')  
  
    plt.ylabel('j values')  
  
    plt.show()  
  
def ACF_PACF_Plot(y, lags):  
  
    acf = sm.tsa.stattools.acf(y, nlags=lags)  
  
    pacf = sm.tsa.stattools.pacf(y, nlags=lags)  
  
    fig = plt.figure()  
  
    plt.subplot(211)  
  
    plt.title('ACF/PACF of the raw data')  
  
    plot_acf(y, ax=plt.gca(), lags=lags)  
  
    plt.subplot(212)  
  
    plot_pacf(y, ax=plt.gca(), lags=lags)  
  
    fig.tight_layout(pad=3)  
  
    plt.show()
```

```

def plotacf(s, k, n):

    # n = len(y)

    # k = number of lags

    t = pd.Series(s)

    acf = acf_cal(t, k)

    acf1 = acf[:::-1][::-1]

    acf2 = acf1 + acf

    x = np.arange(-k, k + 1)

    insig = 1.96 / np.sqrt(len(np.arange(n)))

    plt.stem(x, acf2, markerfmt='o')

    plt.axhspan(-insig, insig, alpha=0.2, facecolor='0.5')

    plt.ylabel('Magnitude')

    plt.xlabel('Lags')

    plt.title('Autocorrelation Plot')

    plt.show()

```

```

# # CLEANING THE DATA

```

```

# In[168]:

```

```
# Checking the amount of nan values in the data
```

```
df.isnull().sum()
```

```
# In[169]:
```

```
# Using forward fill to handle nan values
```

```
df.ffmpeg(axis='rows', inplace=True)
```

```
# In[170]:
```

```
df.isnull().sum()
```

```
# In[171]:
```

```
# Resample the data, because of computational time.
```

```
# Reduces the number of observations of 2075259 to 34589, structure is still retained
```

```
df_resample = df.resample('H').mean()
```

```
# In[172]:
```

```
# Plot of the dependent variable against time
```

```
y = df_resample['Global_active_power']
```

```
plt.figure()
```

```
plt.plot(y,)
```

```
plt.title('Active Power vs Time - Resampled')
```

```
plt.xlabel('Time')
```

```
plt.ylabel('Active Power (kilowatt)')
```

```
plt.xticks(rotation=45)
```

```
plt.show()
```

```
# # STATIONARITY CHECK
```

```
# In[173]:
```



```
# ACF/PACF Plot of the dependent variable
```

```
ACF_PACF_Plot(y, 50)
```

```
# In[174]:
```

```
# Correlation matrix
```

```
plt.figure()
```

```
corr = df_resample.corr()
```

```
ax = sns.heatmap(corr, vmin=-1, vmax=1, center=0, cmap=sns.diverging_palette(150, 275, s=80,  
l=55, n=9), square=True)
```

```
ax.set_xticklabels(ax.get_xticklabels(), rotation=45, horizontalalignment='right')
```

```
plt.title('Correlation Plot')
```

```
plt.tight_layout()
```

```
plt.show()
```

```
# In[175]:
```

```
# Stationary Check on raw data

# the null for ADF is that the time series non-stationary

ADF_Cal(y)


# In[176]:


# the null for KPSS is that the time series is stationary

sm.tsa.stattools.kpss(y, regression='ct')


# In[177]:


# plot of rolling mean and rolling variance of raw data

Rmean = rolling_mean(y)

Rvar = rolling_var(y)

plt.figure()

plt.plot(Rmean)
```

```
plt.title('Rolling Mean - Raw')
```

```
plt.xlabel('t')
```

```
plt.ylabel('y(t)')
```

```
plt.grid()
```

```
plt.show()
```

```
plt.figure()
```

```
plt.plot(Rvar)
```

```
plt.title('Rolling Variance - Raw')
```

```
plt.xlabel('t')
```

```
plt.ylabel('y(t)')
```

```
plt.grid()
```

```
plt.show()
```

```
# Stationary so we can continue
```

```
# # SPLITTING THE DATA
```

```
# In[178]:
```

```
# Splitting the data into Train and Test
```

```
# Split the dataset into train set 80% and test set 20%
```

```
y_train, y_test = train_test_split(df_resample, shuffle=False, test_size=0.2)
```

```
#X =
```

```
df[['Global_reactive_power','Voltage','Global_intensity','Sub_metering_1','Sub_metering_2','Sub_metering_3','Sub_metering_4']]
```

```
#x_train, y_train = train_test_split()
```

```
# # TIME SERIES DECOMPOSITION
```

```
# In[179]:
```

```
# Time Series Decomposition
```

```
ActivePower = df_resample['Global_active_power']
```

```
ActivePower = pd.Series(np.array(df_resample['Global_active_power']),
                        index=pd.date_range('2006-12-16 17:00:00', periods=len(ActivePower)),
                        name='Active Power (kilowatt)')
```

```
# In[180]:
```

```
from statsmodels.tsa.seasonal import STL

STL = STL(ActivePower)

res = STL.fit()

T = res.trend

S = res.seasonal

R = res.resid

# Seasonally adjusted data and plot it vs the original

sadjusted = ActivePower - S

detrended = ActivePower - T

plt.figure(figsize=(8,6))

plt.plot(ActivePower, label= 'Original')

plt.plot(sadjusted, label='Seasonally Adjusted')

plt.title("Original vs Seasonal Adjustment")

plt.xlabel("t")

plt.ylabel("Active Power (kilowatt)")

plt.legend()

plt.show()
```

```

plt.figure(figsize=(8,6))

plt.plot(ActivePower, label= 'Original')

plt.plot(detrended, label='Detrended')

plt.title("Original vs Detrended")

plt.xlabel("t")

plt.ylabel("Active Power (kilowatt)")

plt.legend()

plt.show()

```

```

# In[181]:

```

```

Ft = np.max([0,1 - np.var(R)/np.var(T+R)])

Fs = np.max([0,1 - np.var(R)/np.var(S+R)])

print("The strength of trend for this dataset is ", Ft)

print("The strength of seasonality for this dataset is", Fs)

```

```

# # HOLTS WINTER FORECAST

```

```
# In[182]:
```

```
# Holts winter method
```

```
# use training data to fit model
```

```
model = ets.ExponentialSmoothing(y_train['Global_active_power'], damped_trend= True,  
                                seasonal_periods=12, trend='mul', seasonal='mul').fit()
```

```
# prediction on train set
```

```
HW_train = model.forecast(steps=len(y_train['Global_active_power']))
```

```
HW_train = pd.DataFrame(HW_train,  
columns=['Global_active_power']).set_index(y_train.index)
```

```
# prediction on test set
```

```
HW_test = model.forecast(steps=len(y_test['Global_active_power']))
```

```
HW_test = pd.DataFrame(HW_test, columns=['Global_active_power']).set_index(y_test.index)
```

```
# In[183]:
```

```
# forecast error
```

```
HW_FE = (y_test['Global_active_power'].values).flatten() -  
HW_test['Global_active_power'].values.flatten()  
  
# forecast MSE  
  
HW_MSE = np.round(np.mean(np.square(np.subtract(HW_test['Global_active_power'].values,  
y_test['Global_active_power'].values))), 4)  
  
# In[184]:  
  
# MODEL ASSESSMENT  
  
# In[185]:  
  
print("The mean of the error of the HoltWinter Model is", np.mean(HW_FE))  
print("The variance of error of the HoltWinter Model is :", np.var(HW_FE))  
print("The MSE of the HoltWinter Model is", HW_MSE)  
print("The RMSE of the HoltWinter Model is", np.sqrt(HW_MSE))
```



```
# In[186]:
```

```
plt.figure(figsize=(10, 5))  
  
plt.plot(y_train['Global_active_power'], label='Train')  
  
plt.plot(y_test['Global_active_power'], label='Test')  
  
plt.plot(HW_test, label='Holt Winter Forecast')  
  
plt.title('Holt Winter')  
  
plt.xlabel('t')  
  
plt.ylabel('Kilowatt')  
  
plt.legend()  
  
plt.show()
```

```
# In[187]:
```

```
plotacf(HW_FE, 20, len(HW_FE))
```

```
# In[188]:
```

```
print(sm.stats.acorr_ljungbox(HW_FE, lags=[20], boxpierce=True, return_df=True))
```

```
# In[189]:
```

```
# Residuals are not white, model is not a good fit
```

```
# # FEATURE SELECTION
```

```
# In[303]:
```

```
X =
```

```
df_resample[['Global_reactive_power','Voltage','Global_intensity','Sub_metering_1','Sub_metering_2','Sub_metering_3','Sub_metering_4']]
```

```
X = sm.add_constant(X)
```

```
Y = df_resample[['Global_active_power']]
```

```
X_train, X_test, Y_train, Y_test = train_test_split(X,Y, shuffle = False, test_size = 0.2)
```

```
# OLS Model
```

```
model = sm.OLS(Y_train, X_train).fit()
```

```
print(model.summary())
```

```
# In[304]:
```

```
# Removing Global intensity P-value: 0.810
```

```
X.drop('Global_intensity', axis=1, inplace=True)
```

```
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, shuffle=False, test_size=0.2)
```

```
model = sm.OLS(Y_train, X_train).fit()
```

```
print(model.summary())
```

```
# In[305]:
```

```
# Removing const - P-value:0.486
```

```
X.drop('const', axis=1, inplace=True)
```

```
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, shuffle=False, test_size=0.2)
```

```
model = sm.OLS(Y_train, X_train).fit()
```

```
print(model.summary())
```

```
# # MULTIPLE LINEAR REGRESSION MODEL
```

```
# In[306]:
```

```
model = sm.OLS(Y_train, X_train).fit()
```

```
Tr_pred = model.predict(X_train)
```

```
T_pred = model.predict(X_test)
```

```
print(model.summary())
```

```
# In[307]:
```

```
X = X_train.values
```

```
print("The condition number for X is", LA.cond(X))
```

```
H = np.matmul(X.T,X)
```

```
s,d,v = np.linalg.svd(H)
```

```
print("SingularValues are", d)
```

```
# In[308]:
```

```
plt.figure(figsize=(10, 5))

plt.plot(Y_train, label='Train') # Y_train Values

plt.plot(Y_test, label='Test') # Y_test Values

plt.plot(Tr_pred, label='Train Pred') # Train Predictions

plt.plot(T_pred, label='Test Predictions') # Test Predictions using X_test

plt.legend(loc='best')

plt.xlabel('T')

plt.ylabel('Kilowatt')

plt.title('Train vs Test Predictions - Multiple Linear Regression Model')

plt.show()
```

```
# In[314]:
```

```
H_pred =

model.predict(df_resample.drop(columns=['Global_active_power', 'Global_intensity']))
```

```
# In[316]:
```

```
plt.figure(figsize=(10, 5))  
  
plt.plot(df_resample['Global_active_power'], label='Actual')  
  
plt.plot(H_pred, label='H step Predictions')  
  
plt.legend(loc='best')  
  
plt.xlabel('T')  
  
plt.ylabel('Kilowatt')  
  
plt.title('H step - Multiple Linear Regression Model')  
  
plt.show()
```

```
# In[313]:
```

```
df_resample
```

```
# In[196]:
```

```
#t-test
```

```
print("The p-value of t-test is: ",model.pvalues)
```

```
#f-test
```

```
print("The p-value of f-test is: ",model.f_pvalue)
```

```
# In[197]:
```

```
print("The AIC value of the model is :",model.aic)
```

```
print("The BIC value of the model is :",model.bic)
```

```
print("The R-squared value of the model is: ",model.rsquared)
```

```
print("The Adjusted R-squared value of the model is: ",model.rsquared_adj)
```

```
# In[198]:
```

```
OLS_PE = Y_train.values.flatten() - Tr_pred.values
```

```
OLS_FE = Y_test.values.flatten() - T_pred.values  
OLS_MSE = round(np.square(OLS_FE).mean(),2)  
print("The MSE of the residual is :", OLS_MSE)
```

```
# In[199]:
```

```
print(sm.stats.acorr_ljungbox(OLS_FE, lags=[20], boxpierce=True, return_df=True))
```

```
# In[200]:
```

```
OLS_acf = acf_cal(pd.Series(OLS_FE),20)
```

```
# In[201]:
```

```
OLS_acf2 = acf_cal(pd.Series(OLS_PE), 20)
```



```
# In[202]:
```

```
plotacf(OLS_FE, 20, len(OLS_FE))
```

```
# In[203]:
```

```
plotacf(OLS_PE, 20, len(OLS_PE))
```

```
# In[204]:
```

```
# Q value
```

```
Q_ML_FE = len(OLS_FE)*sum(np.square(OLS_acf[1:]))
```

```
print("The Q value is:", Q_ML_FE)
```

```
# In[205]:
```

```

print("The variance of Prediction error is:", np.var(OLS_PE))

print("The mean of Prediction error is:", np.mean(OLS_PE))

print("The variance of Forecast error is:", np.var(OLS_FE))

print("The mean of Forecast error is:", np.mean(OLS_FE))

```

```

# ##### BASE MODELS #####

```

```

# # AVERAGE METHOD

```

```

# In[206]:

```

```

#y_train1 = y_train.values

```

```

#y_test1 = y_test.values

```

```

h_step = []

```

```

for i in range(len(y_test)):

```

```

    h_step.append(np.mean(y_train['Global_active_power'].values))

```

```
# In[207]:
```

```
AV_FE = y_test['Global_active_power'].values.flatten() - np.array(h_step).flatten()
```

```
AV_MSE = round(np.square(AV_FE).mean(),2)
```

```
print("The MSE of the residual for the Average Method is :", AV_MSE)
```

```
print("The RMSE of the residual for the Average Method is:", np.sqrt(AV_MSE))
```

```
# In[208]:
```

```
print(sm.stats.acorr_ljungbox(AV_FE, lags=[20], boxpierce=True, return_df=True))
```

```
# In[209]:
```

```
AV_acf = acf_cal(pd.Series(AV_FE),20)
```

```
# In[210]:
```

```
plotacf(AV_FE, 20, len(AV_FE))
```

```
# In[211]:
```

```
# Q value
```

```
Q_AV_FE = len(AV_FE)*sum(np.square(AV_acf[1:]))
```

```
print("The Q value is:", Q_AV_FE)
```

```
# In[212]:
```

```
print("The variance of Forecast error is:", np.var(AV_FE))
```

```
print("The mean of Forecast error is:", np.mean(AV_FE))
```

```
# In[213]:
```

```
plt.figure(figsize=(10, 5))

plt.plot(Y_train,label= "Train Data")

plt.plot(Y_test,label= "Test Data")

plt.plot(y_test.index,h_step, label= "Average Method Forecast")

plt.legend(loc='best')

plt.title('Train vs Test Predictions - Average Method Model')

plt.xlabel("t")

plt.ylabel("Kilowatt")

plt.show()
```

```
# # NAIVE METHOD
```

```
# In[214]:
```

```
h_stepN = []

for i in range(len(y_test)):

    h_stepN.append(y_train['Global_active_power'][-1])
```

```
# In[215]:
```

```
N_FE = y_test['Global_active_power'].values.flatten() - np.array(h_stepN).flatten()
```

```
N_MSE = round(np.square(N_FE).mean(),2)
```

```
print("The MSE of the residual for the Average Method is :", N_MSE)
```

```
print("The RMSE of the residual for the Average Method is:", np.sqrt(N_MSE))
```

```
# In[216]:
```

```
print(sm.stats.acorr_ljungbox(N_FE, lags=[20], boxpierce=True, return_df=True))
```

```
# In[217]:
```

```
N_acf = acf_cal(pd.Series(N_FE),20)
```

```
# In[218]:
```

```
plotacf(N_FE, 20, len(N_FE))
```

```
# In[219]:
```

```
# Q value
```

```
Q_N_FE = len(N_FE)*sum(np.square(N_acf[1:]))
```

```
print("The Q value is:", Q_N_FE)
```

```
# In[220]:
```

```
print("The variance of Forecast error is:", np.var(N_FE))
```

```
print("The mean of Forecast error is:", np.mean(N_FE))
```

```
# In[221]:
```

```

plt.figure(figsize=(10, 5))

plt.plot(Y_train,label= "Train Data")

plt.plot(Y_test,label= "Test Data")

plt.plot(y_test.index,h_stepN, label= "Naive Method Forecast")

plt.legend(loc='best')

plt.title('Train vs Test Predictions - Naive Method Model')

plt.xlabel("t")

plt.ylabel("Kilowatt")

plt.show()


# # DRIFT METHOD


# In[222]:


h_stepDR = []

for i in range(1,len(y_test) + 1):

    slope = (y_train['Global_active_power'][-1] - y_train['Global_active_power'][0]) /

    (len(y_train)-1)

```



```
h_stepDR.append(y_train['Global_active_power'][-1]+ i*slope)
```

```
# In[223]:
```

```
DR_FE = y_test['Global_active_power'].values.flatten() - np.array(h_stepDR).flatten()
```

```
DR_MSE = round(np.square(DR_FE).mean(),2)
```

```
print("The MSE of the residual for the Drift Method is :", DR_MSE)
```

```
print("The RMSE of the residual for the Drift Method is:", np.sqrt(DR_MSE))
```

```
# In[224]:
```

```
print(sm.stats.acorr_ljungbox(DR_FE, lags=[20], boxpierce=True, return_df=True))
```

```
# In[225]:
```

```
DR_acf = acf_cal(pd.Series(DR_FE),20)
```

```
# In[226]:
```

```
plotacf(DR_FE, 20, len(DR_FE))
```

```
# In[227]:
```

```
# Q value
```

```
Q_DR_FE = len(DR_FE)*sum(np.square(DR_acf[1:]))
```

```
print("The Q value is:", Q_DR_FE)
```

```
# In[228]:
```

```
print("The variance of Forecast error is:", np.var(DR_FE))
```

```
print("The mean of Forecast error is:", np.mean(DR_FE))
```

```
# In[229]:
```

```
plt.figure(figsize=(10, 5))

plt.plot(Y_train,label= "Train Data")

plt.plot(Y_test,label= "Test Data")

plt.plot(y_test.index,h_stepDR, label= "Drift Method Forecast")

plt.legend(loc='best')

plt.title('Train vs Test Predictions - Drift Method Model')

plt.xlabel("t")

plt.ylabel("Kilowatt")

plt.show()
```

```
# # SES METHOD
```

```
# In[230]:
```

```
fit = SimpleExpSmoothing(np.asarray(y_train['Global_active_power'])).fit(smoothing_level=0.5,
optimized=False)
```

```
h_stepSES = fit.forecast(len(y_test))
```

```
# In[231]:
```

```
SES_FE = y_test['Global_active_power'].values.flatten() - np.array(h_stepSES).flatten()
```

```
SES_MSE = round(np.square(SES_FE).mean(),2)
```

```
print("The MSE of the residual for the Average Method is :", SES_MSE)
```

```
print("The RMSE of the residual for the Average Method is:", np.sqrt(SES_MSE))
```

```
# In[232]:
```

```
print(sm.stats.acorr_ljungbox(SES_FE, lags=[20], boxpierce=True, return_df=True))
```

```
# In[233]:
```

```
SES_acf = acf_cal(pd.Series(SES_FE),20)
```

```
# In[234]:
```

```
plotacf(SSES_FE, 20, len(SSES_FE))
```

```
# In[235]:
```

```
# Q value
```

```
Q_SES_FE = len(SSES_FE)*sum(np.square(SSES_acf[1:]))
```

```
print("The Q value is:", Q_SES_FE)
```

```
# In[236]:
```

```
print("The variance of Forecast error is:", np.var(SSES_FE))
```

```
print("The mean of Forecast error is:", np.mean(SSES_FE))
```

```
# In[237]:
```

```
plt.figure(figsize=(10, 5))

plt.plot(Y_train,label= "Train Data")

plt.plot(Y_test,label= "Test Data")

plt.plot(y_test.index,h_stepSES, label= "SES Forecast")

plt.legend(loc='best')

plt.title('Train vs Test Predictions - SES Method Model')

plt.xlabel("t")

plt.ylabel("Kilowatt")

plt.show()
```

```
# # ARMA MODEL
```

```
# In[238]:
```

```
y_acf = acf_cal(pd.Series(y.values),20)
```

```
# In[239]:
```

```
GPAC_plot(y_acf, 8, 8)
```

```
# In[240]:
```

```
# potentially (2,0) or (1,0) or (2,1)
```

```
# In[242]:
```

```
model = sm.tsa.ARMA(y,(1,0)).fit(trend='nc',disp=0)
```

```
# In[243]:
```

```
print(model.summary())
```

```
# In[244]:
```

```
for i in range(1):
```

```
    print("The AR estimated coefficient a{}".format(i), "is:", model.params[i])
```

```
for i in range(1):
```

```
    print("The confidence interval for estimated coefficient a{}".format(i), "is:", model.conf_int())
```

```
# In[245]:
```

```
print("The standard deviation of the parameter estimates is: ",model.summary().tables[1])
```

```
# In[255]:
```



```
print("The corvariance of estimated parameters is:",model.cov_params())
```

```
# In[246]:
```

```
print("The standard deviation of the parameter estimate is: ",np.square(model.sigma2))
```

```
# In[247]:
```

```
# Prediction with ARMA
```

```
onestep_ARMA = model.predict(start=0, end=len(y_train)-1)
```

```
hstep_ARMA = model.predict(start=len(y_train)-1, end=len(df_resample))
```

```
# In[248]:
```

```
ARMA_PE = y_train['Global_active_power'].values.flatten() - np.array(onestep_ARMA).flatten()
```

```
ARMA_FE = y_test['Global_active_power'].values.flatten() - np.array(hstep_ARMA[2:]).flatten()
```

```
ARMA_MSE = round(np.square(ARMA_FE).mean(),2)

print("The MSE of the residual for the Average Method is :", ARMA_MSE)

print("The RMSE of the residual for the Average Method is:", np.sqrt(ARMA_MSE))
```

```
# In[249]:
```

```
print(sm.stats.acorr_ljungbox(ARMA_FE, lags=[20], boxpierce=True, return_df=True))
```

```
# In[250]:
```

```
ARMA_acf = acf_cal(pd.Series(ARMA_FE),20)
```

```
# In[251]:
```

```
plotacf(ARMA_FE, 20, len(ARMA_FE))
```

```
# In[252]:
```

```
# Q value
```

```
Q_ARMA_FE = len(ARMA_FE)*sum(np.square(ARMA_acf[1:]))
```

```
print("The Q value is:", Q_ARMA_FE)
```

```
# In[253]:
```

```
print("The variance of Prediction error is:", np.var(ARMA_PE))
```

```
print("The mean of Prediction error is:", np.mean(ARMA_PE))
```

```
print("The variance of Forecast error is:", np.var(ARMA_FE))
```

```
print("The mean of Forecast error is:", np.mean(ARMA_FE))
```

```
print("The variance of prediction vs forecast error is:", ARMA_PE.var()/ARMA_FE.var())
```

```
# In[254]:
```

```
plt.figure(figsize=(10, 5))

plt.plot(Y_train,label= "Train Data")

plt.plot(Y_test,label= "Test Data")

plt.plot(y_test.index,hstep_ARMA[2:], label= "ARMA Forecast")

plt.legend(loc='best')

plt.title('Train vs Test Predictions - ARMA Model')

plt.xlabel("t")

plt.ylabel("Kilowatt")

plt.show()
```

```
# In[256]:
```

```
model2 = sm.tsa.ARMA(y,(2,1)).fit(trend='nc',disp=0)
```

```
# In[257]:
```

```
print(model2.summary())
```

```
# In[147]:
```

```
for i in range(2):
```

```
    print("The AR coefficient a{}".format(i), "is:", model2.params[i])
```

```
for i in range(1):
```

```
    print("The MA coefficient a{}".format(i), "is:", model2.params[i + 2])
```

```
# In[258]:
```

```
for i in range(1):
```

```
    print("The confidence interval for estimated coefficient is:", model2.conf_int())
```

```
# In[150]:
```

```
print("The standard deviation of the parameter estimates is: ", np.square(model2.sigma2))
```

```
# In[260]:
```

```
print("the covariance Matrix for the data is", model2.cov_params())
```

```
# In[261]:
```

```
# Prediction with ARMA(2,1)
```

```
onestep_ARMA2 = model2.predict(start=0, end=len(y_train)-1)
```

```
hstep_ARMA2 = model2.predict(start=len(y_train)-1, end=len(df_resample))
```

```
# In[262]:
```

```
ARMA_PE2 = y_train['Global_active_power'].values.flatten() -
```

```
np.array(onestep_ARMA2).flatten()
```

```
ARMA_FE2 = y_test['Global_active_power'].values.flatten() -
```

```
np.array(hstep_ARMA2[2:]).flatten()
```

```
ARMA_MSE2 = round(np.square(ARMA_FE2).mean(),2)

print("The MSE of the residual for the Average Method is :", ARMA_MSE2)

print("The RMSE of the residual for the Average Method is:", np.sqrt(ARMA_MSE2))


# In[153]:


print(sm.stats.acorr_ljungbox(ARMA_FE2, lags=[20], boxpierce=True, return_df=True))


# In[154]:


ARMA_acf2 = acf_cal(pd.Series(ARMA_FE2),20)


# In[155]:


plotacf(ARMA_FE2, 20, len(ARMA_FE2))
```

```
# In[263]:
```

```
# Q value
```

```
Q_ARMA_FE2 = len(ARMA_FE2)*sum(np.square(ARMA_acf2[1:]))
```

```
print("The Q value is:", Q_ARMA_FE2)
```

```
# In[264]:
```

```
print("The variance of Prediction error is:", np.var(ARMA_PE2))
```

```
print("The mean of Prediction error is:", np.mean(ARMA_PE2))
```

```
print("The variance of Forecast error is:", np.var(ARMA_FE2))
```

```
print("The mean of Forecast error is:", np.mean(ARMA_FE2))
```

```
print("The variance of prediction vs forecast error is:", ARMA_PE2.var()/ARMA_FE2.var())
```

```
# In[266]:
```



```

plt.figure(figsize=(10, 5))

plt.plot(Y_train,label= "Train Data")

plt.plot(Y_test,label= "Test Data")

plt.plot(y_test.index,hstep_ARMA2[2:], label= "ARMA Forecast")

plt.legend(loc='best')

plt.title('Train vs Test Predictions - ARMA Model(2,1)')

plt.xlabel("t")

plt.ylabel("Kilowatt")

plt.show()

```

```

# # LSTM

```

```

# In[ ]:

```

```

# code for LSTM does not belong to me, it was adopted from Susan Li fro towardsdatascience
# https://towardsdatascience.com/time-series-analysis-visualization-forecasting-with-lstm-
77a905180eba

```

```

# In[112]:

```

```
get_ipython().system('pip install tensorflow')
```

```
# In[121]:
```

```
import keras
```

```
from keras.models import Sequential
```

```
from keras.layers import Dense
```

```
from keras.layers import LSTM
```

```
from keras.layers import Dropout
```

```
from keras.layers import *
```

```
from sklearn.preprocessing import MinMaxScaler
```

```
from sklearn.metrics import mean_squared_error
```

```
from sklearn.metrics import mean_absolute_error
```

```
from keras.callbacks import EarlyStopping
```

```
# In[275]:
```

```

data = df_resample.Global_active_power.values #numpy.ndarray

data = data.astype('float32')

data = np.reshape(data, (-1, 1))

scaler = MinMaxScaler(feature_range=(0, 1))

data = scaler.fit_transform(data)

train_size = int(len(data) * 0.80)

test_size = len(data) - train_size

train, test = data[0:train_size:], data[train_size:len(data),:]

```

```
# In[276]:
```

```
# convert an array of values into a dataset matrix
```

```
def create_dataset(dataset, look_back=1):
```

```
    X, Y = [], []
```

```
    for i in range(len(dataset)-look_back-1):
```

```
        a = dataset[i:(i+look_back), 0]
```

```
        X.append(a)
```

```
        Y.append(dataset[i + look_back, 0])
```

```
    return np.array(X), np.array(Y)
```

```
# In[277]:
```

```
# reshape into X=t and Y=t+1
```

```
look_back = 30
```

```
X_train, Y_train = create_dataset(train, look_back)
```

```
X_test, Y_test = create_dataset(test, look_back)
```

```
# In[278]:
```

```
X_train.shape
```

```
# In[279]:
```

```
Y_train.shape
```

```
# In[280]:
```

```
# reshape input to be [samples, time steps, features]
```

```
X_train = np.reshape(X_train, (X_train.shape[0], 1, X_train.shape[1]))
```

```
X_test = np.reshape(X_test, (X_test.shape[0], 1, X_test.shape[1]))
```

```
# In[281]:
```

```
model = Sequential()
```

```
model.add(LSTM(100, input_shape=(X_train.shape[1], X_train.shape[2])))
```

```
model.add(Dropout(0.2))
```

```
model.add(Dense(1))
```

```
model.compile(loss='mean_squared_error', optimizer='adam')
```

```
history = model.fit(X_train, Y_train, epochs=20, batch_size=70, validation_data=(X_test, Y_test),
```

```
                    callbacks=[EarlyStopping(monitor='val_loss', patience=10)], verbose=1,
```

```
                    shuffle=False)
```

```
# Training Phase
```

```
model.summary()
```

```
# In[282]:
```

```
# make predictions
```

```
train_predict = model.predict(X_train)
```

```
test_predict = model.predict(X_test)
```

```
# invert predictions
```

```
train_predict = scaler.inverse_transform(train_predict)
```

```
Y_train = scaler.inverse_transform([Y_train])
```

```
test_predict = scaler.inverse_transform(test_predict)
```

```
Y_test = scaler.inverse_transform([Y_test])
```

```
print('Train Mean Absolute Error:', mean_absolute_error(Y_train[0], train_predict[:,0]))
```

```
print('Train Root Mean Squared Error:', np.sqrt(mean_squared_error(Y_train[0],  
train_predict[:,0])))
```

```
print('Test Mean Absolute Error:', mean_absolute_error(Y_test[0], test_predict[:,0]))
```

```
print('Test Root Mean Squared Error:', np.sqrt(mean_squared_error(Y_test[0],  
test_predict[:,0])))
```

```
# In[283]:
```

```
plt.figure(figsize=(8,4))

plt.plot(history.history['loss'], label='Train Loss')

plt.plot(history.history['val_loss'], label='Test Loss')

plt.title('model loss')

plt.ylabel('loss')

plt.xlabel('epochs')

plt.legend(loc='upper right')

plt.show();
```

```
# In[285]:
```

```
aa=[x for x in range(200)]

plt.figure(figsize=(10,5))

plt.plot(Y_test[0], marker='.', label="Test")

plt.plot(test_predict[:,0], 'r', label="Test Pred")
```

```
#plt.plot(Y_train[0], marker='.', label="Train")

#plt.plot(train_predict[:,0], 'g', label="Train Pred")

# plt.tick_params(left=False, labelleft=True) #remove ticks

plt.tight_layout()

sns.despine(top=True)

plt.subplots_adjust(left=0.07)

plt.ylabel('Global_active_power(KW)', size=15)

plt.xlabel('Time step', size=15)

plt.legend(fontsize=15)

plt.title('LSTM Forecast')

plt.show();
```

```
# In[273]:
```

```
len(y_train['Global_active_power'].values.flatten())
```

```
# In[294]:
```



```
LSTM_PE = np.subtract(Y_train[0], train_predict[:,0])  
  
LSTM_PMSE = np.sqrt(mean_squared_error(Y_train[0], train_predict[:,0]))  
  
LSTM_FE = np.subtract(Y_test[0], test_predict[:,0])  
  
LSTM_MSE = np.sqrt(mean_squared_error(Y_test[0], test_predict[:,0]))
```

```
# In[297]:
```

```
LSTM_acf = acf_cal(pd.Series(LSTM_FE),20)
```

```
# In[298]:
```

```
# Q value
```

```
Q_LSTM_FE = len(LSTM_FE)*sum(np.square(LSTM_acf[1:]))  
  
print("The Q value is:", Q_LSTM_FE)
```

```
# In[299]:
```

```
plotacf(LSTM_FE, 20, len(LSTM_FE))
```

```
# In[300]:
```

```
print(sm.stats.acorr_ljungbox(LSTM_FE, lags=[20], boxpierce=True, return_df=True))
```

```
# In[ ]:
```